Identifying the distributional sources of children's early vocabulary

Anonymous CogSci submission

Abstract

Children's early vocabulary learning must to a large extent be driven by the prevalence of words: they can't learn a word if they haven't heard it. Indeed, previous research has found that higher word frequency is a good predictor of earlier learning. However, despite considerable overlap, word frequency distributions also vary significantly by source: child-directed speech, books, and television have distinct profiles. Children receive a mixture of these different frequency distributions, and the ratios of the mixture may be predictive of their early word learning. The goal of this paper is to better understand the shared and unique variance in these sources of input–in both English and French–and to evaluate how predictive these input frequencies are of children's early word learning.

Keywords: early language learning; CDI; vocabulary development; word frequency distributions.

Introduction

Previous studies have shown that frequency matters for children's word learning (for a review, see Ambridge, Kidd, Rowland, & Theakston, 2015), and have observed an association between word frequency in children's language environments and age of acquisition (Goodman, Dale, & Li, 2008). However, input word frequency varies significantly depending on the context. For instance, word frequency in books is not the same as frequency in conversational speech (Dawson, Hsiao, Wei Ming Tan, Banerji, & Nation, 2021; Montag, Jones, & Smith, 2015). Some differences between frequency distributions are intuitive: "mommy" is quite frequent in child-directed speech (2,260 tokens per million; TPM), yet not so common in children's books (10 TPM), and even more rare in books meant for all ages (2 TPM). But other differences are less intuitive: "of" is frequent in books meant for all ages (41,630 TPM), and while still frequent in child-directed speech (5,900 TPM), relatively less so as compared to children's books (20,400 TPM). In general, speech–both directed to children, and to adults–contains relatively fewer function words, and tends to score lower on measures of lexical diversity than books (Dawson et al., 2021).

Different language input sources, such as the use of child-directed register or book reading time, can lead to variance in input speech heard by child during their everyday lives. This input variance has often been interpreted as a function of the families' socioeconomic status (SES; Rowe, 2018). Importantly, this variance has been found to relate to children's language development and to be predictive of aspects of word learning (Hoff, 2003).

The primary goal of this paper is to examine shared and unique variance in word frequency across different sources of input, ranging from children's books and movies to child-directed speech and adult-directed books, movies, and speech, which we accomplish via principle components analysis (PCA). After characterizing the structure of the principle components of frequency from different sources in both English and French, we investigate how well these components predict English- and French-learning children's early word learning, using aggregate MacArthur-Bates Communicative Development Inventories (CDI) data from Wordbank (Frank, Braginsky, Yurovsky, & Marchman, 2017). Finally, we examine how well the frequency components predict individual children's word learning in combination with their mother's education, which may be related to how much children are read to at home.

Method

Datasets

Child-directed Speech. Utterances of child-directed speech (ChS) were extracted from the CHILDES corpus (MacWhinney, 2000), a collection of transcripts of dyadic interactions between caregivers and children of ages ranging from 0 to 12 years ($M = 2.9$ years). After cleaning, the CHILDES corpus yielded a total of 5.521096×10^6 tokens across 38779 word types. The French CHILDES corpus yielded a total of 4.7492515×10^5 tokens across 14310 word types.

Child-directed books (ChdB). We used a sample of 98 children's books from Project Gutenberg's open-source database of books that has been used in prior machine learning research on language comprehension (Hill, Bordes, Chopra, & Weston, 2015). These books were published between 1820 and 1922, but include such well-known titles as *The Legend of Sleep Hollow*. After cleaning, this children's book corpus totals 4.674202×10^6 tokens across 42666 word types. For French, we used 130 popular children stories openly accessible in parenting websites (`https://fr.hellokids.com/`) and X books from Gutemberg. After cleaning, this children's book corpus totals 9.7590699×10^5 tokens across word types.

Child-directed Media (ChM). Transcripts from television shows (e.g., from PBS Kids and Nickelodeon) and movies (e.g., *Beauty and the Beast*), including 1,078 movies and 4,309 TV episodes were taken from Charlesworth, Yang, Mann, Kurdi, & Banaji (2021) (available here: `https://osf.io/kqux5/`. After cleaning, this children's media corpus totals 6.759247×10^6 tokens across 84333 word types. For French, we also used openly accessible movie subtitles (`https://www.subsynchro.com/`), X number of films. After cleaning, this children's media corpus totals 9.3599787×10^5 tokens across 1402 word types.

Adult-directed Speech (AdS). Adult-directed speech was obtained from the Switchboard-1 Telephone Speech Corpus (Godfrey & Holliman, 1993), a corpus of transcripts from dyadic telephone conversations in which 543 adult speakers were assigned to discuss a randomly-assigned topic. After cleaning, the adult-directed speech corpus (AdS) yielded a total of 3.104328×10^6 tokens across 27536 word types. For French, adult-directed speech was obtained from the TCOF corpus of adult speech transcripts collected during the 80s-90s for research (André & Canut, 2010), from the CLAPI corpus (Balthasar & Bert, 2005) and the CFPP corpus (Branca-Rosoff, Fleury, Lefeuvre, & Pires, 2012). The adult-directed speech corpus (AdS) yielded a total of 8.8567847×10^5 tokens across 883 word types.

Adult-directed Books (AdB). The adult-directed book corpus (AdB) is comprised of 3.5143617×10^7 tokens across 827414 word types. For French, the corpus is comprised of books taken from the 1999 Association de Bibliophiles Universels, an open-source database of french books. It is comprised of 9.1965581×10^5 tokens across 3849 word types.

Adult-directed Media (AdM). The adult-directed media corpus (AdM) is comprised of 6.166909×10^6 tokens across 62876 word types. For French, the corpus is comprised of 9.3916931×10^5 tokens across 1710 word types, using openly accessible movie subtitles (`https://www.subsynchro.com/`), X number of films.

Merging the Corpora

Children's early word learning data is drawn from the CDIs (Fenson et al., 2007). CDIs are parental reports on their children's lexical development, proven to be reliable indicators of a child's language. CDIs are survey instruments, where parents mark whether their child (age ranges 8-15 and 16-30 months old) understands or produces particular words out of a list of several hundred words.

All word frequencies were normalized to number of tokens per million (TPM). We focus our analysis on the 674 words from the English CDI that we were able to find in at least some of the corpora, and 470 words from the French CDI (for French, words were matched to related words in corpora via a stemmer). For any CDI words that failed to appear in a given corpus, we replaced the missing word's frequency with a normalized count of 10 TPM, or the minimum normalized frequency for that distribution, whichever was smaller.

Results

Cross-corpus Frequency Correlations

Table 1 shows the cross-corpus word frequency correlations for the 674 CDI words.

Table 1: Correlation table of word frequency distributions for English CDI words.

	ChS	ChM	ChB	AdS	AdM	AdB
ChS	1.00	0.83	0.73	0.83	0.75	0.62
ChM	0.83	1.00	0.94	0.87	0.97	0.89
ChB	0.73	0.94	1.00	0.87	0.94	0.91
AdS	0.83	0.87	0.87	1.00	0.83	0.74
AdM	0.75	0.97	0.94	0.83	1.00	0.96
AdB	0.62	0.89	0.91	0.74	0.96	1.00

Table 2: Correlation table of word frequency distributions for French CDI words.

	ChS	ChM	ChB	AdS	AdM	AdB
ChS	1.00	0.89	0.70	0.88	0.84	0.67
ChM	0.89	1.00	0.77	0.93	0.97	0.77
ChB	0.70	0.77	1.00	0.86	0.81	0.97
AdS	0.88	0.93	0.86	1.00	0.92	0.85
AdM	0.84	0.97	0.81	0.92	1.00	0.84
AdB	0.67	0.77	0.97	0.85	0.84	1.00

Age of Acquisition Regression

Despite the strong correlations of word frequency across these different corpora, we will attempt a simple regression predicting each CDI word's mean Age of Acquisition (AoA) – the mean age (in months) at which 50% of children are expected to know a given word (Braginsky, Yurovsky, Marchman, & Frank, 2019). We also include the number of letters as a predictor (Nletters) to help control for the overall difficulty of each word. Table 2 below displays the estimated coefficients of this regression, showing significant contributions of all child-directed distributions (speech, books, and media), and of adult speech frequencies, as well as the number of letters. However, the variance inflation factor (VIF) values for all of the frequency distributions are all $\gg 1$ (and many > 5 or > 10), indicating that these variables show strong multicollinearity which may compromise the reliability of the regression results. Thus, we turned to principal components analysis (PCA) to disentangle these correlated distributions and to understand their interrelations.

Principal Components of Frequency

We used PCA to examine the principal components of the six log-scaled word frequency distributions ([adult- vs. child-

	Beta	SE	t-val	p-val	VIF
(Intercept)	27.39	0.90	30.44	0.00	
log(ChS)	-2.16	0.17	-12.37	0.00	4.56
log(ChM)	-0.52	0.23	-2.24	0.03	9.46
log(ChB)	1.42	0.17	8.18	0.00	6.64
log(AdS)	0.59	0.12	5.07	0.00	6.34
log(AdM)	0.18	0.18	1.00	0.32	12.63
log(AdB)	0.18	0.17	1.04	0.30	11.96
Nletters	0.40	0.09	4.71	0.00	1.45

Table 3: Coefficients for English corpora frequencies predicting AoA.

	Beta	SE	t-val	p-val	VIF
(Intercept)	27.34	0.57	48.25	0.00	
log(ChS)	-1.61	0.14	-11.66	0.00	3.26
log(ChM)	0.07	0.13	0.52	0.60	3.78
log(ChB)	0.07	0.13	0.52	0.60	4.30
log(AdS)	0.42	0.13	3.14	0.00	3.96
log(AdM)	0.76	0.13	5.66	0.00	3.87
log(AdB)	0.22	0.13	1.70	0.09	4.16
Nletters	0.01	0.06	0.14	0.89	1.28

Table 4: Coefficients for French corpora frequencies predicting AoA.

directed] x [speech, books, media]). Table 3 shows the standard deviation (Std Dev) and proportion of variance explained, both individually (Prop Var) and cumulatively (Cum Prop Var) by the principal components (PC1-PC6). PC1 already explains the bulk the variance (89%), and PC2-PC4 each only capture an additional 2-4% of the variance. In total, the first four components captured >98% of the variance.

Table 5: Importance of components from English PCA.

	PC1	PC2	PC3	PC4	PC5	PC6
Std Dev	5.10	1.05	0.97	0.82	0.57	0.44
Prop Var	0.89	0.04	0.03	0.02	0.01	0.01
Cum Prop Var	0.89	0.93	0.96	0.98	0.99	1.00

Tables 3 and 4 show the eigenvectors of the principal components (PC1-PC6) in relation to the original six frequency distributions for English and French, respectively. Tables 5 and 6 show the correlation of each frequency distribution with the CDI items' loadings on each principal component for English and French, respectively.

For English, it is clear that PC1 captures overall frequency, but loads more strongly on the adult distributions (-.47 to -.52). PC2 (3.8% of variance) mostly captures child-directed speech (0.65), especially differentiating it from adult-directed books and media (-.47, -.45). PC3 (3.1% of variance) captures adult-directed speech (.79), particularly from child-

Table 6: Importance of components from French PCA.

	PC1	PC2	PC3	PC4	PC5	PC6
Std Dev	2.90	0.92	0.69	0.62	0.59	0.56
Prop Var	0.78	0.08	0.04	0.04	0.03	0.03
Cum Prop Var	0.78	0.86	0.90	0.94	0.97	1.00

directed distributions. PC4 (2.2% of variance) captures the similarity of child- and adult-directed media (.40, .48), distinguishing them from child- and adult-directed books (-.63, -.44). PC5 (1.1% of variance) mostly captures child-directed books (0.53), distinguishing it from child-directed speech (-.60) and adult-directed books (-.50). PC6 (<1% of variance) captures child-directed media (.73), distinguishing it in particular from adult-directed media (-.52). In summary, the principal components align surprisingly well with particular dimensions of the English frequency distributions: PC1 with overall adult-directed frequency, PC2 with child-directed speech, PC3 with adult-directed speech, PC4 with media vs. books, PC5 with child-directed books vs. speech, and PC6 with adult- vs. child-directed media. Although even PC5 and PC6 seem interpretable, given limited space and the fact that they account for very little variance (∼ 1%) we will focus on PC1-PC4 for the rest of our analyses.

For French, we see a similar pattern of findings, although the order of the most important principal components is not quite the same. As seen in Table 6, the first PC for the French corpora captures less of the variance than for the English corpora, and remaining PCs capture more – even PC5 and PC6 each capture 3%, compared to 1% in the English PCA.

Table 7: Principal components' rotation in the original coordinate system.

	PC1	PC2	PC3	PC4	PC5	PC6
ChM	-0.31	0.15	-0.35	0.40	0.28	0.73
ChB	-0.34	0.25	-0.32	-0.63	0.53	-0.21
ChS	-0.26	0.65	-0.30	0.12	-0.60	-0.22
AdM	-0.47	-0.45	-0.24	0.48	0.13	-0.53
AdB	-0.49	-0.47	-0.01	-0.44	-0.50	0.32
AdS	-0.52	0.27	0.79	0.10	0.14	-0.01

PCA-based Prediction of Age of Acquisition

Figure 1 shows the loadings of English CDI items on PC1-PC4 vs. the average age of acquisition (AoA; in months).

Regularized Regression on PCA Loadings

Next we used the CDI items' PCA loadings in lieu of the frequency distributions to predict AoA. First, we did a L1-regularized (i.e., LASSO) regression predicting AoA with all of the principal components, to test which of the PCs should

Table 8: Principal components' rotation in the original coordinate system. (French)

	PC1	PC2	PC3	PC4	PC5	PC6
ChM	-0.41	0.32	-0.13	0.81	-0.23	0.05
ChB	-0.43	-0.57	-0.25	-0.08	-0.03	0.65
ChS	-0.36	0.11	-0.74	-0.25	0.23	-0.44
AdM	-0.40	0.40	0.37	-0.12	0.69	0.24
AdB	-0.43	-0.53	0.45	0.10	0.03	-0.57
AdS	-0.41	0.35	0.20	-0.50	-0.65	0.03

be included in the regression with lexical class. We used cross-validation to find the λ value (penalty for outsized coefficients) that minimized mean-squared error of the test set, resulting in λ = 0.006 for English, a small value which will result in coefficients that would be quite close to those obtained in ordinary least squares regression (i.e., when λ = 0). Using this best-fitting λ, this model had $R^2 = 0.32$ in English, and the estimated coefficients were non-zero for all of the principal components. For English, there were negative coefficients for PC1 (β = −0.17), PC2 (β = −1.13), PC4 (β = −1.32), and PC6 (β = −0.34), indicating that higher values on these PCs predict earlier AoAs – especially for PC2 and PC4, which mostly capture child-directed speech (PC2) and media (PC4). There were positive coefficients for PC3 (β = 0.88) and PC5 (β = 1.95), indicating that higher values of these PCs predict later AoAs, which is reasonable, as PC3 is associated with adult-directed speech, and PC5 captures child-directed books. This is still reasonable, since even child-directed books contain more function words than speech, which tend to be later-learned. Moreover, children receive a relatively small fraction of their daily input from books – perhaps 20 minutes/day, in contrast with several hours of child-directed speech and overheard adult-directed speech.

For French, this model had $R^2 = 0.12$, and the estimated coefficients were non-zero for all of the principal components. There were negative coefficients for PC1 (β = −0.05) and PC5 (β = −0.13), indicating that higher values on these PCs predict earlier AoAs. There were positive coefficients for PC2 (β = 0.14), PC3 (β = 1.63), PC4 (β = 0.18) and PC6 (β = 0.82) indicating that higher values of these PCs predict later AoAs.

Given that past research has found that lexical class strongly modulates influences of word frequency, we next examined the interaction of lexical class (LC) with PC1 - PC6. To determine if the inclusion of all PCs was justified, we ran a series of ANOVAs building up from PC1 to PC6–in decreasing order of the variance they accounted for in the PCA. Thus, the R syntax for the sequence of regressions was `AoA~PC1*LC`, `AoA~(PC1+PC2)*LC`, ..., `AoA~(PC1+PC2+PC3+PC4+PC5+PC6)*LC`, with noun as the baseline LC. The more complex model was always signifi-

cantly preferred, including up to the inclusion of PC6 ($R^2 = .584$), although for ease of presentation we only show the results of the regression with up to PC4. Table 9 shows the results of this regression, which yielded $R^2 = .538$. PC1 and PC3 had significant positive coefficients, while PC2 and PC4 had significant negative coefficients. PC1 had significant interactions with all levels of lexical class, with negative coefficients. PC2 had significant interactions with .. (positive), and with .. (negative). Interactions of PC3 and PC4 with LC were not significant. [better to just show a ggeffects graph?]

	Beta	SE	t-val	p-val
(Intercept)	22.11	0.23	97.32	0.00
PC1	0.52	0.05	10.13	0.00
PC2	-1.97	0.20	-9.97	0.00
PC3	0.93	0.16	5.72	0.00
PC4	-0.84	0.20	-4.21	0.00
LCadj	3.43	0.50	6.92	0.00
LCfunc	7.82	0.66	11.93	0.00
LCother	2.35	0.39	6.03	0.00
LCverb	3.75	0.42	8.83	0.00
PC1:LCadj	-0.51	0.12	-4.32	0.00
PC1:LCfunc	-0.38	0.09	-4.39	0.00
PC1:LCother	-0.47	0.09	-5.47	0.00
PC1:LCverb	-0.44	0.11	-4.15	0.00
PC2:LCadj	0.93	0.51	1.83	0.07
PC2:LCfunc	0.95	0.28	3.37	0.00
PC2:LCother	-0.82	0.37	-2.25	0.02
PC2:LCverb	1.00	0.46	2.16	0.03
PC3:LCadj	-0.56	0.58	-0.96	0.34
PC3:LCfunc	-0.59	0.33	-1.78	0.07
PC3:LCother	0.34	0.35	0.95	0.34
PC3:LCverb	0.20	0.40	0.52	0.60
PC4:LCadj	-0.26	0.68	-0.38	0.71
PC4:LCfunc	-0.35	0.41	-0.86	0.39
PC4:LCother	-0.12	0.39	-0.30	0.76
PC4:LCverb	0.75	0.51	1.47	0.14

Table 9: Regression predicting English CDI AoAs with PCs and lexical class.

Combining distributions with demographic data

Past research has found that young children from higher-SES households tend to have larger vocabulary. Parents with higher-SES tend to also report reading more to their young children than parents with lower SES. Together, this suggests that the vocabulary composition of children from higher-SES households may be better predicted by the word frequencies seen in child-directed books, rather than those from child-directed speech. To test this idea, we did an exploratory logistic regression using the first four PCs to predict the number of children in Wordbank who produce or don't produce each item, along with interactions of mother's education and children's age. American English data from Wordbank contained 2,776 CDI:WS administrations with mother's

education (coded: -1 for no more than secondary education (N=547), 0 for some/all college (N=1483), and 1 for at least some graduate school (N=746)). There were significant main effects of age, mother's education, and PC1-PC4 (all $p < .001$). There were significant interactions of age with mother's education ($p < .001$), PC2 ($p < .001$), and PC4 ($p = .009$). There were significant interactions of mother's education with PC1-PC3 (all $p < .001$). [how to show these effects?]

	Beta	SE	t-val	p-val
(Intercept)	-5.33	0.05	-117.54	0.00
age	0.21	0.00	204.85	0.00
MotherEdCollege	-1.54	0.03	-53.74	0.00
MotherEdGraduate	-1.76	0.03	-55.35	0.00
PC1	0.07	0.01	7.39	0.00
PC2	0.31	0.04	7.53	0.00
PC3	-0.29	0.05	-6.08	0.00
PC4	0.15	0.05	2.81	0.00
age:MotherEdCollege	0.07	0.00	59.18	0.00
age:MotherEdGraduate	0.09	0.00	67.84	0.00
age:PC1	-0.00	0.00	-7.94	0.00
age:PC2	-0.01	0.00	-6.89	0.00
age:PC3	0.00	0.00	3.92	0.00
age:PC4	0.00	0.00	3.03	0.00
MotherEdCollege:PC1	0.04	0.01	7.68	0.00
MotherEdGraduate:PC1	0.06	0.01	8.75	0.00
MotherEdCollege:PC2	0.10	0.03	3.62	0.00
MotherEdGraduate:PC2	-0.03	0.03	-1.02	0.31
MotherEdCollege:PC3	-0.09	0.03	-3.03	0.00
MotherEdGraduate:PC3	-0.02	0.03	-0.64	0.52
MotherEdCollege:PC4	0.02	0.04	0.52	0.60
MotherEdGraduate:PC4	-0.03	0.04	-0.88	0.38
age:MotherEdCollege:PC1	-0.00	0.00	-5.10	0.00
age:MotherEdGraduate:PC1	-0.00	0.00	-6.98	0.00
age:MotherEdCollege:PC2	-0.00	0.00	-1.88	0.06
age:MotherEdGraduate:PC2	0.00	0.00	2.70	0.01
age:MotherEdCollege:PC3	0.00	0.00	1.07	0.28
age:MotherEdGraduate:PC3	-0.00	0.00	-1.63	0.10
age:MotherEdCollege:PC4	-0.00	0.00	-0.25	0.81
age:MotherEdGraduate:PC4	0.00	0.00	0.38	0.71

Table 10: Regression coefficients for predicting English CDI items using PCs and mother's education.

Discussion

We set out to investigate the sources of linguistic input that children may experience using word frequency distributions garnered from child-directed and adult-directed corpora of speech, books, and media (TV and movies). In both English and French corpora, we found the principal components (PCs) of these distributions, and described how these PCs capture variation both in adult- vs. child-directedness, as well as between modalities (e.g., books and speech). Moreover, in both English and French we found that multiple components are predictive of children's age of acquisition of words from the CDI.

(Although a reason for the differences between the English and French findings may be that French corpora were composed by several smaller ones, due to the lack of big accessible corpora online, and the slight different use for the corpora e.g. small stories in French, large books in English).

Finally, we used English Wordbank data to examine how well the frequency components combine with mother's education–a measure of SES that has been found in the past to be positively related to early word learning, and to more child-directed reading–to predict children's early word learning. This analysis revealed significant contributions of multiple PCs, as well as the interaction of mother's education with PC1 (overall frequency), PC2 (child-directed speech), and PC3 (adult-directed speech), but not of the child-directed books component (PC4). A target for future research is to predict individual children's learning of particular words using these principal components, in combination with parent-reported measures of how much time their child spends daily receiving input from each of the input sources ([adult- vs. child-directed] x [books, media, and speech]).

In conclusion, despite the overwhelming similarity of word frequency distributions from different sources, we have shown that these distributions show systematic variation–at least in English and in French, which can predict some of the variation in children's early word learning. By better understanding the similarity and differences between word frequencies children experience in different contexts, future research in this vein holds the promise to predict individual differences in children's early word learning on the basis of their daily routines.

Acknowledgements

References

Ambridge, B., Kidd, E., Rowland, C. F., & Theakston, A. L. (2015). The ubiquity of frequency effects in first language acquisition. *Journal of Child Language*, 42(2), 239–273.

André, V., & Canut, E. (2010). Mise à disposition de corpus oraux interactifs: Le projet tcof (traitement de corpus oraux en français). *Pratiques. Linguistique, Littérature, Didactique*, (147-148), 35–51.

Balthasar, L., & Bert, M. (2005). La plateforme Corpus de langues parlées en interaction(CLAPI). Historique, état des lieux, perspectives. *Lidil. Revue de Linguistique et de Didactique Des Langues*, (31), 13–33.

Braginsky, M., Yurovsky, D., Marchman, V. A., & Frank, M. C. (2019). Consistency and variability in children's word learning across languages. *Open Mind*, 3, 52–67.

Branca-Rosoff, S., Fleury, S., Lefeuvre, F., & Pires, M. (2012). Discours sur la ville. Présentation du corpus de français parlé parisien des années 2000 (cfpp2000). *Article En Ligne, Http://Cfpp2000. Univparis3. Fr/Articles. Html*.

Charlesworth, T. E., Yang, V., Mann, T. C., Kurdi, B., & Banaji, M. R. (2021). Gender stereotypes in natural language: Word embeddings show robust consistency across child and adult language corpora of more than 65 million words. *Psychological Science*, *32*(2), 218–240.

Dawson, N., Hsiao, Y., Wei Ming Tan, A., Banerji, N., & Nation, K. (2021). Features of lexical richness in children's books: Comparisons with child-directed speech. *Language Development Research*.

Fenson, L., Marchman, V. A., Thal, D. J., Dale, P. S., Reznick, J. S., & Bates, E. (2007). *MacArthur-Bates Communicative Development Inventories: User's guide and technical manual (2nd ed.)*. Baltimore, MD: Brookes.

Frank, M. C., Braginsky, M., Yurovsky, D., & Marchman, V. A. (2017). Wordbank: An open repository for developmental vocabulary data. *Journal of Child Language*, *44*(3), 677–694.

Goodman, J. C., Dale, P. S., & Li, P. (2008). Does frequency count? Parental input and the acquisition of vocabulary. *Journal of Child Language*, *35*(3), 515–531.

Hill, F., Bordes, A., Chopra, S., & Weston, J. (2015). The goldilocks principle: Reading children's books with explicit memory representations. *arXiv Preprint arXiv:1511.02301*.

Hoff, E. (2003). The specificity of environmental influence: Socioeconomic status affects early vocabulary development via maternal speech. *Child Development*, *74*(5), 1368–1378.

MacWhinney, B. (2000). *The childes project: Tools for analyzing talk. Transcription format and programs* (Vol. 1). Psychology Press.

Montag, J. L., Jones, M. N., & Smith, L. B. (2015). The words children hear: Picture books and the statistics for language learning. *Psychological Science*, *26*(9), 1489–1496.

Rowe, M. L. (2018). Understanding socioeconomic differences in parents' speech to children. *Child Development Perspectives*, *12*(2), 122–127.
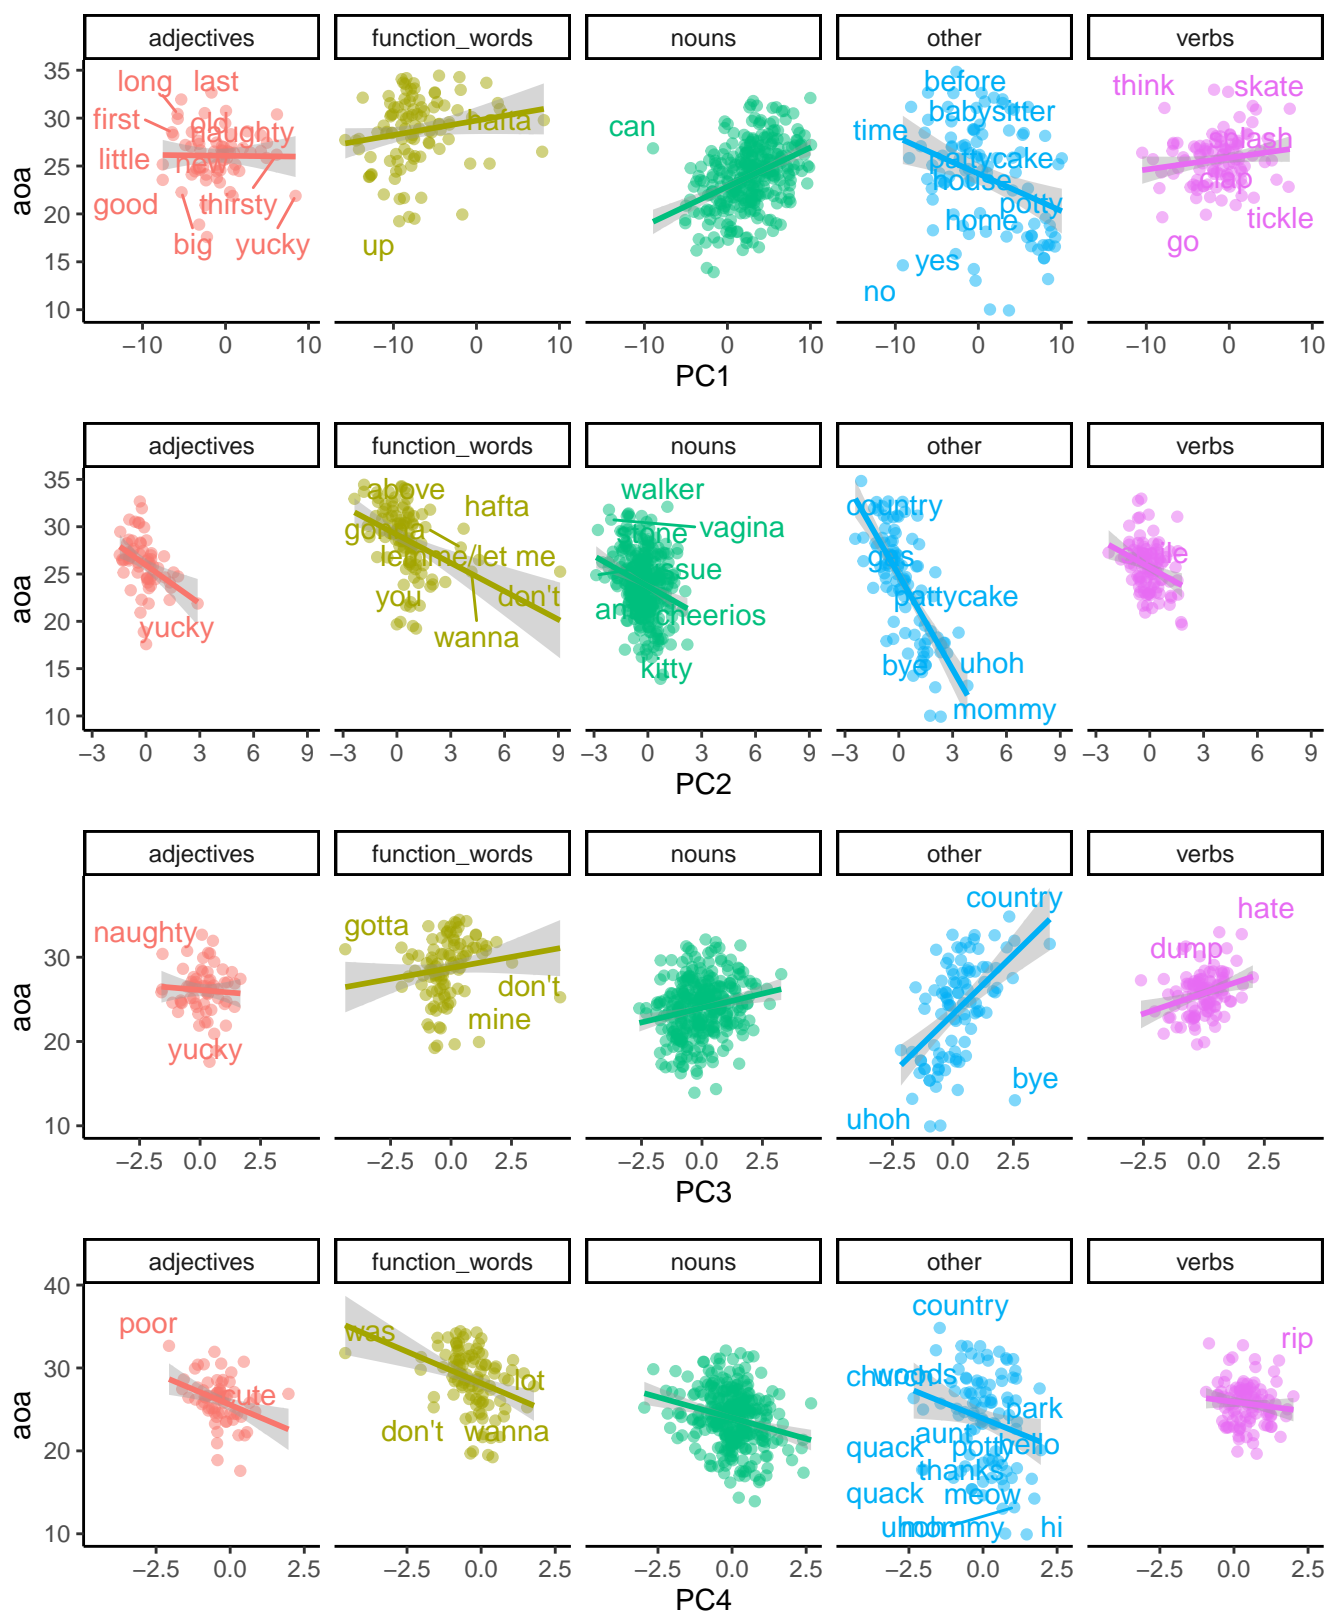
Figure 1: Principal components vs. age of acquisition, by lexical class.